



What's New in Apache Spark 2.3?

Xiao Li

Feb 8, 2018



About Me

- Software Engineer at Databricks
- Apache Spark Committer and PMC Member
- Previously, IBM Master Inventor
- Spark SQL, Database Replication, Information Integration
- Ph.D. in University of Florida
- Github: [gatorsmile](#)



Simplify Big Data and AI with Databricks

Increases Data Science Productivity by 5x

Eliminates Disparate Tools with Optimized Spark

Removes Devops & Infrastructure Complexity

DATABRICKS COLLABORATIVE NOTEBOOKS
Explore Data → Train Models → Serve Models

DATABRICKS RUNTIME

Reliability | APACHE | Performance

DATABRICKS SERVERLESS



Accelerates & Simplifies Data Prep for Analytics

Databricks Enterprise Security



IoT / STREAMING DATA



CLOUD STORAGE



DATA WAREHOUSES



HADOOP STORAGE

Databricks Customers Across Industries

Financial Services



BLACKROCK



JPMORGAN CHASE & CO.



Healthcare & Pharma



HUMAN LONGEVITY,
INC.



REGENERON

Media & Entertainment

NBCUniversal
VIACOM



BETHESDA GAMES STUDIOS

zynga TiVO

LIVE NATION



Mc Graw Hill Education



edmunds

Data & Analytics Services

Cox AUTOMOTIVE™



THOMSON REUTERS



edmunds

Technology

AUTODESK



cisco

adobe

GoPro

NetApp

Public Sector



U.S. Citizenship and Immigration Services

AARP

BLACKSKY
YOUR WORLD NOW

CMS
CENTERS FOR MEDICARE & MEDICAID SERVICES

noblis

DigitalGlobe

databricks

Retail & CPG

Red Bull

ShopRite

Groupon

DOLLAR SHAVE CLUB

flipp

Westfield

overstock.com®

Consumer Services

Expedia

Hotels.com

mapquest

Blue Apron

glassdoor

myfitnesspal

Flipboard

OpenTable

Medium

AIMIA

RADIUS

INNERACTIVE

Marketing & AdTech

MediaMath

SOCIAL CODE

TUNE

LoyaltyOne

eyevue

AIMIA

RADIUS

INNERACTIVE

Energy & Industrial IoT

Shell

DNV·GL

Quby

SIEMENS

ConocoPhillips

databricks

This Talk

Databricks Runtime 4.0



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



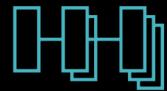
Spark on
Kubernetes



Databricks
Delta

OSS Apache Spark 2.3

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream
Join



UDF
Enhancements



Image Reader



Native ORC
Support



Stable
Codegen



Various SQL
Features

Over 1300 issues resolved!

This Talk



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



Databricks
Delta

Structured Streaming

stream processing on Spark SQL engine

fast, scalable, fault-tolerant

rich, unified, high level APIs

deal with *complex data* and *complex workloads*

rich ecosystem of data sources

integrate with many *storage systems*

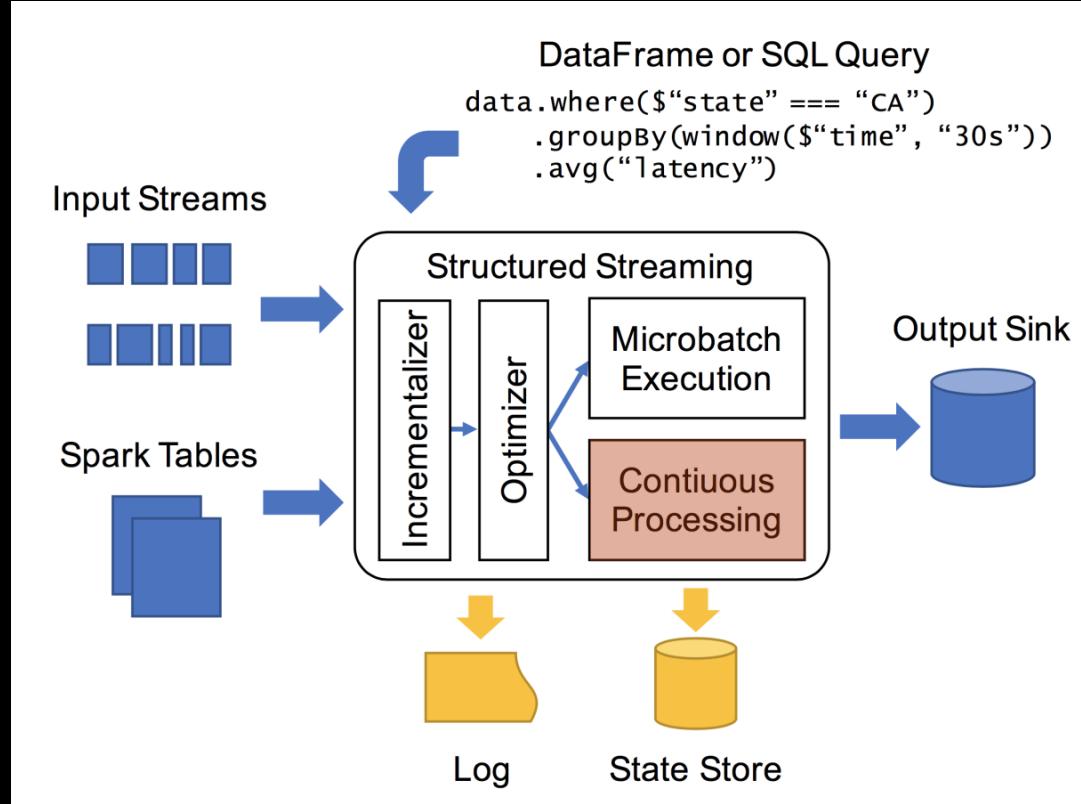
Structured Streaming

Introduced in Spark 2.0

Among Databricks customers:

- 10X more usage than DStream
- Processed 100+ trillion records in production

Structured Streaming



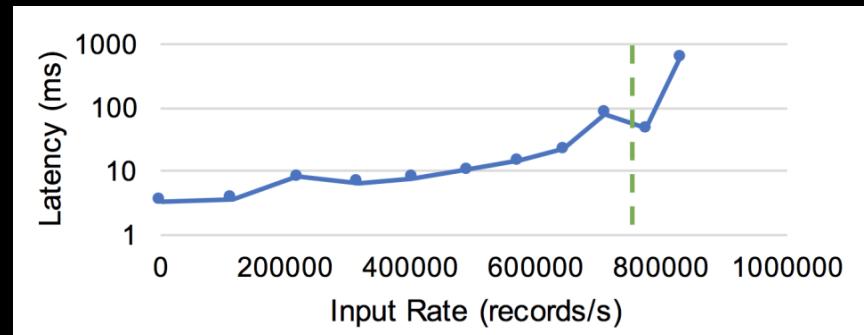
Continuous Processing Execution Mode

Micro-batch Processing (since 2.0 release)

- Lower end-to-end latencies of ~100ms
- Exactly-once fault-tolerance guarantees

Continuous Processing (since 2.3 release) [[SPARK-20928](#)]

- A new streaming execution mode (experimental)
- Low (~1 ms) end-to-end latency
- At-least-once guarantees





Continuous
Processing

```
import org.apache.spark.sql.streaming.Trigger

spark
    .readStream
    .format("rate")
    .option("rowsPerSecond", "10")
    .option("")

spark
    .readStream
    .format("kafka")
    .option("kafka.bootstrap.servers", "host1:port1,host2:port2")
    .option("subscribe", "topic1")
    .load()
    .selectExpr("CAST(key AS STRING)", "CAST(value AS STRING)")
    .writeStream
    .format("kafka")
    .option("kafka.bootstrap.servers", "host1:port1,host2:port2")
    .option("topic", "topic1")
    .trigger(Trigger.Continuous("1 second"))
    .start()
```

The only change you need!



Continuous
Processing

Continuous Processing

Supported Operations:

- Map-like Dataset operations
 - projections
 - selections
- All SQL functions
 - Except `current_timestamp()`,
`current_date()` and
aggregation functions

Supported Sources:

- Kafka source
- Rate source

Supported Sinks:

- Kafka sink
- Memory sink
- Console sink

This Talk



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



Databricks
Delta



ML on
Streaming

ML on Streaming

Model transformation/prediction on batch and streaming data with unified API.

After fitting a model or Pipeline, you can deploy it in a streaming job.

```
val streamOutput = transformer.transform(streamDF)
```

Demo

Notebook:

<https://databricks-prod-cloudfront.cloud.databricks.com/public/4027ec902e239c93eaaa8714f173bcfc/1212468499849361/649180143457189/7151759854724023/latest.html>



Image Support in Spark

Image
Reader

Spark Image data source [SPARK-21866](#) :

- Defined a standard API in Spark for reading images into DataFrames
- Deep learning frameworks can rely on this.

```
val df = ImageSchema.readImages("/data/images")
```

This Talk



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



Databricks
Delta

PySpark

Introduced in Spark 0.7 (~2013); became first class citizen in the DataFrame API in Spark 1.3 (~2015)

Much slower than Scala/Java with user-defined functions (UDF), due to serialization & Python interpreter

Note: Most PyData tooling, e.g. Pandas, numpy, are written in C++

PySpark Performance

Fast data serialization and execution using vectorized formats

[\[SPARK-22216\]](#) [\[SPARK-21187\]](#)

- Conversion from/to Pandas

`df.toPandas()`

`createDataFrame(pandas_df)`

- Pandas/Vectorized UDFs: UDF using Pandas to process data
 - Scalar Pandas UDFs
 - Grouped Map Pandas UDFs



PySpark Performance



Blog "Introducing Vectorized UDFs for PySpark" <http://dbricks.co/2rMwmW0>

Pandas UDF

Scalar Pandas UDFs:

- Used with functions such as *select* and *withColumn*.
- The Python function should take *pandas.Series* as inputs and return a *pandas.Series* of the same length.

```
import pandas as pd
from pyspark.sql.functions import col, pandas_udf
from pyspark.sql.types import LongType

x = pd.Series([1, 2, 3])
# Create a Spark DataFrame, 'spark' is an existing SparkSession
df = spark.createDataFrame(pd.DataFrame(x, columns=["x"]))

# Declare the function and create the UDF
def multiply_func(a, b):
    return a * b
multiply = pandas_udf(multiply_func, returnType=LongType())

df.select(multiply(col("x"), col("x"))).show()
```

Pandas UDF

Grouped Map Pandas UDFs:

- **Split-apply-combine**
- A Python function that defines the computation for each group.
- The output schema

```
from pyspark.sql.functions import pandas_udf, PandasUDFType

df = spark.createDataFrame(
    [(1, 1.0), (1, 2.0), (2, 3.0), (2, 5.0), (2, 10.0)],
    ("id", "v"))

@pandas_udf("id long, v double", PandasUDFType.GROUPED_MAP)
def subtract_mean(pdf):
    # pdf is a pandas.DataFrame
    v = pdf.v
    return pdf.assign(v=v - v.mean())

df.groupby("id").apply(subtract_mean).show()
```

Demo

Notebook:

<https://databricks-prod-cloudfront.cloud.databricks.com/public/4027ec902e239c93eaaa8714f173bcfc/1212468499849361/421914239386310/7151759854724023/latest.html>

This Talk



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



Databricks
Delta

● Kubernetes
Search term

● Hadoop
Search term

+ Add comparison

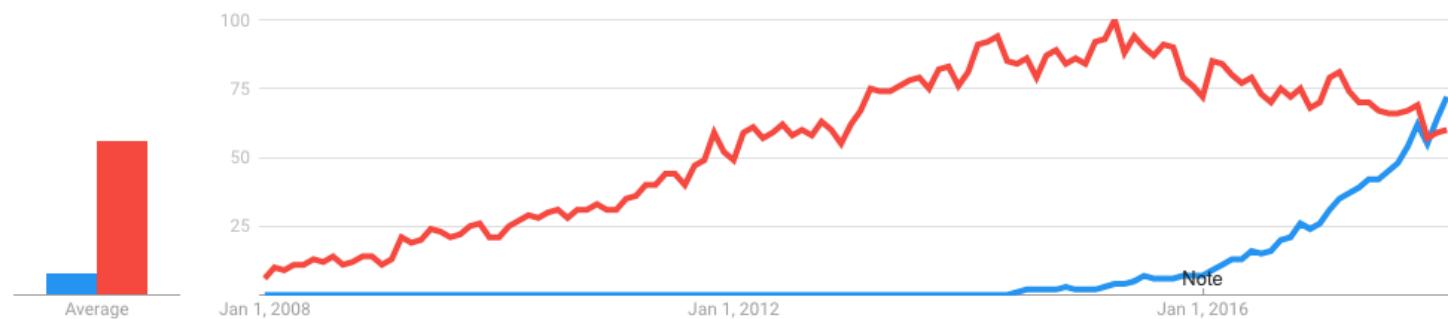
Worldwide ▾

1/1/08 - 2/7/18 ▾

All categories ▾

Web Search ▾

Interest over time



Native Spark App in K8S

- New Spark scheduler backend
- Driver runs in a Kubernetes pod created by the submission client and creates pods that run the executors in response to requests from the Spark scheduler. [[K8S-34377](#)] [[SPARK-18278](#)]
- Make direct use of Kubernetes clusters for multi-tenancy and sharing through [Namespaces](#) and [Quotas](#), as well as administrative features such as [Pluggable Authorization](#), and [Logging](#).



Spark on Kubernetes

Supported:

- Supports Kubernetes 1.6 and up
- Supports cluster mode only
- Static resource allocation only
- Supports Java and Scala applications
- Can use container-local and remote dependencies that are downloadable

In roadmap (2.4):

- Client mode
- Dynamic resource allocation + external shuffle service
- Python and R support
- Submission client local dependencies + Resource staging server (RSS)
- Non-secured and Kerberized HDFS access (injection of Hadoop configuration)

This Talk



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



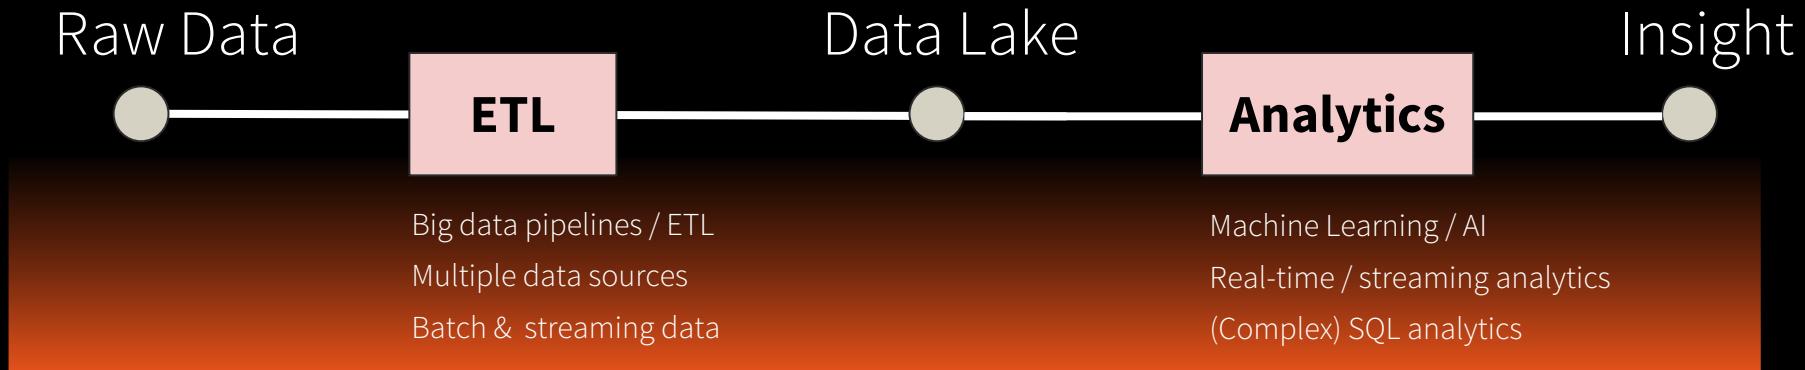
Databricks
Delta

What is Databricks Delta?

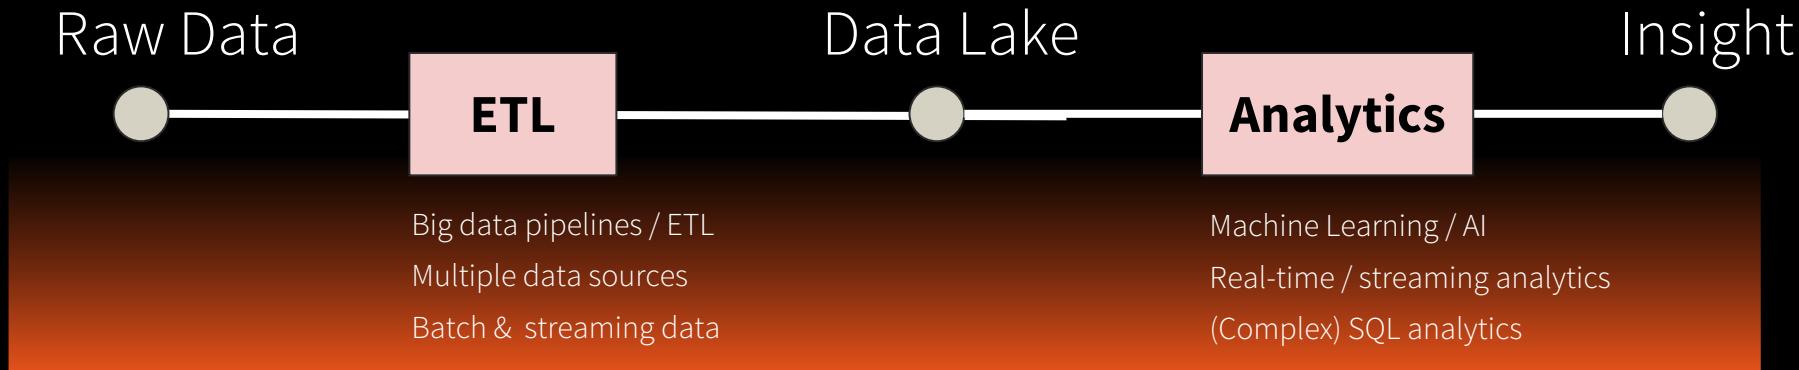
Delta is a data management capability that brings data reliability and performance optimizations to the cloud data lake.

Key Capabilities - Transactions & Indexing

Standard data pipelines with Spark



Yet challenges still remain



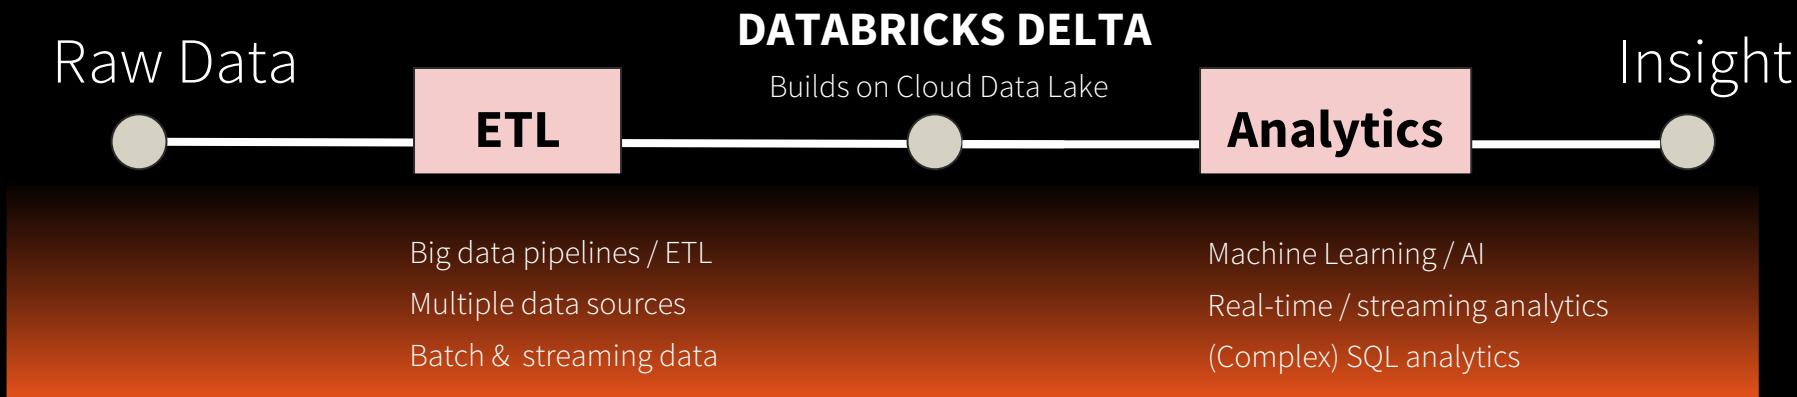
Reliability and Complexity

- Data corruption issues and broken pipelines
- Complex workarounds - tedious scheduling and multiple jobs/staging tables
- Many use cases require updates to existing data - not supported by Spark / Data lakes

Reliability & Performance Problems

- Performance degradation at scale for advanced analytics
- Stale and unreliable data slows analytic decisions

Databricks Delta address these challenges



Reliability & Automation

Transactions guarantees eliminates complexity
Schema enforcement to ensure clean data
Upserts/Updates/Deletes to manage data changes
Seamlessly support streaming and batch

Performance & Reliability

Automatic indexing & caching
Fresh data for advanced analytics
Automated performance tuning

Databricks Delta Under the Hood

SCALE

- Use cloud object storage (e.g. AWS S3)
- Decouple compute & storage

RELIABILITY

- ACID compliant & data upserts
- Schema evolution

PERFORMANCE

- Intelligent partitioning & compaction
- Data skipping & indexing

LOW LATENCY

- Real-time streaming ingest
- Structured streaming



Continuous
Processing



Streaming ML
+
Image Reader



PySpark
Performance



Spark on
Kubernetes



Databricks
Delta

Try these on Databricks Runtime 4.0 beta today!

Community edition:
<https://community.cloud.databricks.com/>



Thank you

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



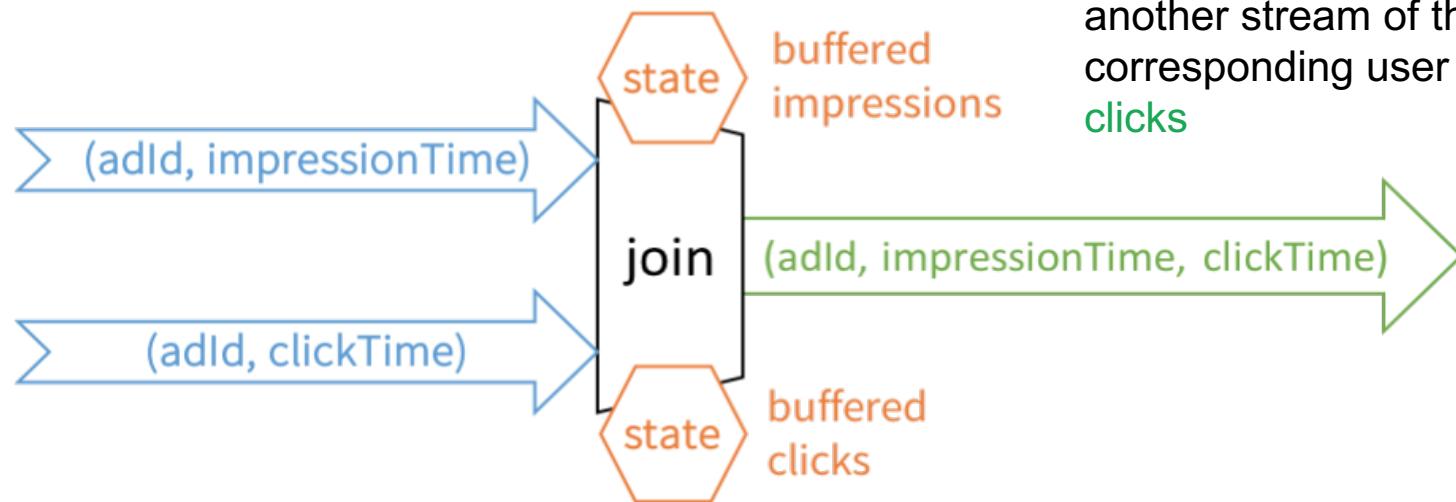
Various SQL Features



Stream-stream Joins

Stream-stream
Join

Example: Ad Monetization





Inner Join + Time constraints + Watermarks

Stream-stream
Join

Time constraints

- Impressions can be 2 hours late
- Clicks can be 3 hours late
- A click can occur within 1 hour after the corresponding impression

```
val impressionsWithWatermark = impressions
    .withWatermark("impressionTime", "2 hours")
```

```
val clicksWithWatermark = clicks
    .withWatermark("clickTime", "3 hours")
```

```
impressionsWithWatermark.join(
    clicksWithWatermark,
    expr"""
        clickAdId = impressionAdId AND
        clickTime >= impressionTime AND
        clickTime <= impressionTime + interval 1 hour
    """
)

```

Range Join



Stream-stream Joins

Stream-stream
Join

Inner	Supported, optionally specify watermark on both sides + time constraints for state cleanup
Left	Conditionally supported, must specify watermark on right + time constraints for correct results, optionally specify watermark on left for all state cleanup
Right	Conditionally supported, must specify watermark on left + time constraints for correct results, optionally specify watermark on right for all state cleanup
Full	Not supported.

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features



Data Source
API V2

What's Wrong With V1?

- **Leak** upper level API in the data source (DataFrame/SQLContext)
- **Hard** to extend the Data Source API for more optimizations
- **Zero** transaction guarantee in the write APIs



Data Source
API V2

Design Goals

- Java friendly (written in Java).
- No dependency on upper level APIs (DataFrame/RDD/...).
- Easy to extend, can add new optimizations while keeping backward compatibility.
- Can report physical information like size, partition, etc.
- Streaming source/sink support.
- A flexible and powerful, transactional write API.
- No change to end users.



Data Source
API V2

Features in Spark 2.3

- Support both row-based scan and columnar scan.
- Column pruning and filter pushdown.
- Can report basic statistics and data partitioning.
- Transactional write API.
- Streaming source and sink support for both micro-batch and continuous mode.
- Spark 2.4 + : more enhancements

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features



UDF Enhancements

UDF Enhancement

- [SPARK-19285] Implement UDF0 (SQL UDF that has 0 arguments)
- [SPARK-22945] Add java UDF APIs in the functions object
- [SPARK-21499] Support creating SQL function for Spark UDAF(UserDefinedAggregateFunction)
- [SPARK-20586][SPARK-20416][SPARK-20668] Annotate UDF with Name, Nullability and Determinism

```
val foo = udf(() => Math.random()).asNondeterministic()
testData.select(foo())
```



Java UDF and UDAF

UDF
Enhancements

- Register Java UDF and UDAF as a SQL function and use them in PySpark.

```
# class MyDoubleAvg extends UserDefinedAggregateFunction
spark.udf.registerJavaUDAF("javaUDAF", "test.org.apache.spark.sql.MyDoubleAvg")
df = spark.createDataFrame([(1, "a"),(2, "b"), (3, "a")],["id", "name"])
df.registerTempTable("df")
spark.sql("SELECT name, javaUDAF(id) as avg from df group by name").collect()

# class JavaStringLength implements UDF1<String, Integer>
spark.udf.registerJavaFunction(
    "javaStringLength", "test.org.apache.spark.sql.JavaStringLength", "integer")
spark.sql("SELECT javaStringLength('test')").collect()
```

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features



Stable
Codegen

Stable Codegen

- [SPARK-22510] [SPARK-22692] Stabilize the codegen framework to avoid hitting the 64KB JVM bytecode limit on the Java method and Java compiler constant pool limit.
- [SPARK-21871] Turn off whole-stage codegen when the bytecode of the generated Java function is larger than **spark.sqlcodegen.hugeMethodLimit**. The limit of method bytecode for JIT optimization on HotSpot is 8K.

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features

Vectorized ORC Reader

- [SPARK-20682] Add new ORCFileFormat based on ORC 1.4.1.
spark.sql.orc.impl = `native` (default) / `hive` (Hive 1.2.1)
- [SPARK-16060] Vectorized ORC reader
spark.sql.orc.enableVectorizedReader = `true` (default) / `false`
- Enable filter pushdown for ORC files by default

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features

Performance

- [\[SPARK-21975\]](#) Histogram support in cost-based optimizer
- Enhancements in rule-based optimizer and planner
 - [\[SPARK-22489\]](#) [\[SPARK-22916\]](#) [\[SPARK-22895\]](#) [\[SPARK-20758\]](#)
 - [\[SPARK-22266\]](#) [\[SPARK-19122\]](#) [\[SPARK-22662\]](#) [\[SPARK-21652\]](#)
(e.g., constant propagation)
- [\[SPARK-20331\]](#) Broaden support for Hive partition pruning predicate pushdown. (e.g. date = 20161011 or date = 20161014)
- [\[SPARK-20822\]](#) Vectorized reader for table cache

API

- Improved ANSI SQL compliance and Dataset/DataFrame APIs
- More built-in functions [[SPARK-20746](#)]
- Better Hive compatibility
 - [[SPARK-20236](#)] Support Dynamic Partition Overwrite
 - [[SPARK-17729](#)] Enable Creating Hive Bucketed Tables
 - [[SPARK-4131](#)] Support INSERT OVERWRITE DIRECTORY

Major Features on Spark 2.3



Continuous Processing



Data Source API V2



Spark on Kubernetes



PySpark Performance



ML on Streaming



History Server V2



Stream-stream Join



UDF Enhancements



Image Reader



Native ORC Support



Stable Codegen



Various SQL Features



History
Server V2

History Server Using K-V Store

Stateless and non-scalable History Server V1:

- Requires parsing the event logs (that means, so slow!)
- Requires storing app lists and UI in the memory (and then OOM!)

[SPARK-18085] **K-V store-based History Server V2:**

- Store app lists and UI in a persistent K-V store (LevelDB)
 - **spark.history.store.path** – once specified, LevelDB is being used; otherwise, in-memory KV-store (still stateless like V1)
 - V2 still can read the event log wrote by V1



Thank you